

Notes - 12/9

Friday, December 09, 2011

9:53 AM

52

[In 41/1753 today at 10:15am]

Agenda:

- Update on engineering and hiring
 - Progress toward 12/16 drop
 - jQuery
 - Test
- Make calls on current spec/implementation divergences:
 - Private
 - Statically-extensible interfaces and extern classes
 - Don't need an 'extension' modifier - we don't do this for modules or elsewhere
 - Don't do extern class extension
 - Others?
- Question: Should we consider including some targeted ES6 features? (destructuring per
- Design work:
 - Mixins and inheritance
 - Generics
 - 'import'
 - Module 'code-loading'
 - Inference for lambdas
 - Abstract classes
- Anything else?

Generics:

View is generic method is bundle of overloads

JSConf: May 1-3

Inheritance:

- Issue with 'this' in
- Suggestion: Allow extending interfaces
 - Extending a class means normal prototype inheritance
 - You can extend multiple classes, the first is the spine
 - Extending an interface means assume it will be there at runtime
 - Extending any interfaces makes you non-instantiable

```
interface Ifoo {  
  foo(): string;  
}
```

```
class Base {  
  foo()=> "Base";  
}
```

```
class A extends Base { // regular extensions  
  
}
```

```
class B extends IFoo { //  
  foo() => super.foo() + "A";  
}
```

ime and type 'super' wrt it

```
class Bus extends Vehicle() {}
class Bicyc extends Vehicle() {}
```

```
class Bicycle extends Vehicle() {}
```

ModuleSourceElement:

GetAccessorPropertyDeclaration

SetAccessorPropertyDeclaration

...

GetAccessorPropertyDeclaration

SetAccessorPropertyDeclaration

GetAccessorPropertyDeclaration:

[illegible]

```
,
```

```
class Fancy extends D2, A {  
  
}
```

x with A(...) // requires x be an Ifoo

- Or does it need to be "extend x with A(...)";
- Idempotent - do we need a marker

Decision: Go for the extends/implements simplification with extending in

Implementation: For now, remove 'requires'.

e's what I
dn't
of using
t
code.

Re-rooting modules

Two common Node.js modules are Net and HTTP. Both are global objects. This is hard to write in Strada today:

```
// This is wrong because there are not  
// Net and HTTP, and these module names  
// type instances returned from require  
module Net {  
  class Server() {}  
}  
module HTTP {  
  class Server() {}  
}  
  
// To deal with the above, we just provide  
// shapes of Net and HTTP. However, to  
// within them for scoping. So we eliminate  
// name conflicts  
interface INet {  
  Server: { new(): IServer }  
}  
interface IHTTP {  
  Server: { new(): IServer }  
}
```

interfaces. No abstracts.

both of these expose a Server constructor

```
let global objects named  
es cannot be used to  
re('http')
```

```
provide interfaces for the module  
then we can't nest the interfaces  
evate to the top level, and get
```

```
property get Identifier ( ) ReturnTypeAnnotationopt
    FunctionImplementation;
```

SetAccessorPropertyDeclaration:

```
property set Identifier ( NamedParameter ) ReturnTypeAnnotationopt
    FunctionImplementation;
```

Type Checking:

Classes should allow properties which use ES5 accessor syntax. Type-wise, they represent a single property. If both present, the parameter type of the setter must be assignable to the return type of the getter. If a getter is present, the type of the property is the return type of the getter. If a getter is not present, the type of the property is the parameter type of the setter. The getter must have zero parameters, and the setter must have one parameter.

Code Generation:

The combination of get and set accessors for a name is combined into a single statement in the output, as a call to:

```
Object.defineProperty(/*class-prototype*/,
    Identifier,
    { get: function Identifier() { /
    *GetAccessorPropertyDeclaration-FunctionImplementation*/ },
      set: function Identifier(NamedParameter) {
    *SetAccessorPropertyDeclaration-FunctionImplementation*/ },
    enumerable: true,
    configurable: true });
```

Luke

Yesterday we decided that

- a) It probably makes sense to think of generic methods as an infinite bundle of overloads, and
- b) We should try addressing calls to generic methods *without* the presence of other overloads.

```

    }
    interface IServer { // Is this the ne
    }

```

We really do want to use modules here, but those modules are at the top level. Instead, they can be requested and rooted to support this using ES6 syntax.

```

//consumer.str
module net from 'net'
module http from 'http'

var server: http.Server = http.createServer()

```

Imports can still be used in conjunction with this, but it's up to the developer to avoid these and keep the module namespace clean. These name conflicts can be dealt with statically in the code to individual conflicting names – though ES6 is stricter.

```

//consumer.str
module net from 'net'
module http from 'http'
import * from http

var server: Server = createServer();

```

This is still not quite right though. An attempt to accommodate this will not be there at runtime. This is because of the static value hierarchy. What we really want is a way to define the module. Notable, with classes and interfaces, we want to define the “type-only” portion and how to define the “type-plus-instance” concept. Interfaces can partly play this role for types. We’re missing something like a ‘moduleface’:

```

moduleface Net {
    class Server() {
    }
}
moduleface HTTP {
    class Server(): Net.Server() {
    }
}

```


`net.Server` or `http.Server`?

modules are not going to be guaranteed to exist at
based on what the user wants. I believe we can

`Server();`

it will potentially reintroduce name conflicts. It is up
es separate if there are name conflicts. (In principle,
he same way as C#, lazily rejecting static references
er than this currently).

ess the value `net.Server` will statically appear to
cause ‘module’ conflates the static type hierarchy
t here is the ability to import only the “type side” of
e have the pair of both how to describe the “type-
nce” portion. For modules, we do not have the
ole also – but cannot themselves contain

declarations first.

So here are some examples looking at calls to generic methods from the viewpoint of “overload bundles”. I am making one more limiting assumption and not thinking about lambdas passed to functions yet – instead when methods take function parameters I am restricting to passing functions that already have a declared type. That way we can worry about dealing with the contravariance of function types first.

I do want to non-humbly tease that there is *awesome insight* towards the end, but you should read in sequence 😊. After, I deeply appreciate any thoughts you have.

Mads

One covariant T in parameters:

```
m<T>(t:T): T;  
var r = m(7);
```

Here it seems clear that `r` should get the type `number`. Out of the infinite bundle it seems that the type `number` is “best”. There could be a couple of reasons *why* it is best, though. First of all we need to decide whether to allow “reverse assignability” (assigned less specific to more specific types) for arguments to methods. If we don’t, then only overloads with the type `number` and its supertype can apply in the first place. If we do, then all the overloads with enum types, plus the `null` and `undefined` types (which are subtypes of `number`) *also* apply.

We probably want to say that overloads requiring reverse assignability of arguments either *don’t* apply at all, or only apply if there aren’t any others. If that is the case, then the set of applicable overloads here is those where `T` is `number` or a subtype.

There are still a couple of potential reasons we could employ why `number` should be the return type. Either because the `number` overload is *best* by some rule (e.g. subtyping on parameter types), or because `number` is the winner, or summary, or intersection of all the *return* types of all the applicable overloads.

One contravariant T in parameters:

```
m<T>(f:(t:T)=>string): T;  
g: (x:number)=>string;  
var r = m(g);
```

$$\left\{ \begin{aligned} & \text{--- } x_0 \\ & \vdots \\ & \dots \end{aligned} \right\}$$

```
Net.Server // error! There is no value
```

```
var server : Net.Server = require('http').createServer(handler)
```

This has several benefits: (1) it allows types to be pro

Recommendation: I don't think any of these options moduleface does seem like a hole we will want to try adding additional basic tools for making typing more more easily than building a holistic module-based co abstract over all possible module systems in current

References and command line

Let's say I want to compile `stradac-node.str`, and it contains `node.str`. However, I do not want to re-emit the `.js` files generated, or may be in folders I don't even have write access to. In `stradac-node.str`, there is no need to code-gen the `node.str`.

This command line will try to compile them all to ind

```
stradac stradac-node.str strada.str ..
```

This command line will concatenate them all into one

```
stradac stradac-node.str strada.str ..
```

What we want in principle is:

```
stradac stradac-node.str /r:strada.str
```

That is, the auxiliary strada files are references, not standalone value environments for type checking, but not involving

Now, in principle we want most of these sorts of con
within source code, so that I do not need a project fi

ue Net, similar to interfaces

```
cp').createServer(); // okay!
```

provided completely independently of structure

s is yet good enough to pursue, but the notion of
y and address in some form. I like the direction of
flexible. This can be adapted to other use cases
de-loading system into Strada which tries to
use.

depends on types from strada.str and ../lib
files for the latter two files – they may already be
ite permissions to. Either way, for compilation of
thers.

individual output files:

```
../lib/node.str
```

e, which is also not desired:

```
../lib/node.str -out stradac-node.js
```

```
.. /r:../lib/node.str
```

sources. They are used for populating the type and
ved in code generation at all.

command line options to be available to be set from
to drive tooling

This example is a little weird. I take a function of something and return that something? Surely a bit contrived? What kind of function would do that? We can make it more realistic by returning a function of `T` instead of a `T`:

```
m<T>(f:(t:T)=>string): (t:T)=>bool;
g: (x:number)=>string = ...;
var r = m(g);
```

Now you can imagine that `m` e.g. composes the incoming function with some post-processing to return a `string` that results in a `bool`. The question is, what is the type of `r`? Clearly we'd like it to be `(x:number)=>bool`, but let's see what we can get from rules like what we applied above. First, the applicable overloads of `m` are those where the parameter type for `f` is `(x:number)=>string` or a supertype thereof. Wait, what is a supertype of a function type? It is one where the return type is more general (but they are all `string` here so that doesn't matter) *or the parameter types are specific!!!* For instance, `m<{}>` is *not* applicable, but `m<E>` for some enum type `E` *is*.

So the set of applicable overloads is mighty different than before, in that the `T`'s are an infinite set of *subtypes* rather than a finite set of *supertypes*. But it doesn't matter much. The most specific parameter type (by subtyping) is still the one where `T` is `number`. Remember, those other candidates were applicable precisely because their parameter types were supertypes of `(x:number)=>string`. The one is a subtype of all the others.

Similarly if we look at return types: If we pick the most specific of the return types it will be the one. If we combine them all with intersection or whatever, it will be the right one. So, somewhat surprisingly perhaps, nothing new here compared to the covariant case above.

Two covariant T's in parameters:

```
m<T>(a: T, b: T): T;
var r = m(bicycle, bus);
```

Imagine a `choose` method which returns one of its two arguments based on the phases of the moon. The applicable overloads are those where `T` is a supertype of `Bicycle` and also of `Bus`. This amounts to finding the most specific common supertype; let's call it `Vehicle` (though it may not be defined by that name anywhere in the program). Clearly the overload with `Vehicle` is the "best" by all the definitions above, so the only thing new here is that we have to accept synthesizing best common supertypes (which isn't particularly hard to do in a structural system).

within source code, so that I do not need a project file

```
//stradac-node.str  
#reference 'strada.str'  
#reference '..\lib\node.str'  
  
var compiler = new Tools.StradaCompiler  
// ...
```

This replaces the current comment driven model. The “reference”, not a dynamic import. The # also helps of executable code.

Now, ideally – I could just write the following, and the source, just like other tools will need to:

```
stradac stradac-node.str
```

This relies on treating the “#reference” like one of the

Recommendation: Adopt /r and #reference as static dynamic code loading for now.

ie to drive tooling.

er());

he use of '#reference' keeps the notion of being a
separate this as a metadirective instead of a piece

he compiler resolves the references implicit in the

he "/r" command line options.

-only reference mechanisms, separate this from

Two contravariant T's in parameters:

```
m<T>(f:(t:T)=>string, g:(t:T)=>number): (t:T)=>bool;  
h: (x:Bicycle)=>string;  
i: (x:Bus)=>number;  
var r = m(h,i);
```

OK, let's say that this method returns a function that returns true if the length of the string returned by `f` applied to its argument `t` is greater than the number returned by `g` applied to `t`. (if that just makes your brain to mush, don't worry about it – just saying that the signature is plausible). The applicable overloads are those where `T` is a subtype of `Bicycle` and also of `Bus`. This amounts to finding the most general common subtype; say `Buscycle`. While it is rather unlikely that any buscycles actually exist, it is no harder for the compiler to produce that type than it is to compute a most specific common subtype, as we did above.

While applicable overloads exist using *subtypes* of `Buscycle` (`Buscyclerries` and `Buscyclanes` for instance), the one with `Buscycle` is clearly the *least* useless, and indeed the one rules like above would choose. Note that it is no accident that the result here is rather useless. We are trying to combine functions that only work for bicycles with functions that only work for busses. There shouldn't be a whole lot of common ground there. So it is no failure of the system to produce this useless result, a failure of the scenario.

More interesting would be if we had used `Vehicle` instead of `Bus` in the example. Here the outcome would be (using all suggested approaches) the overload based on `Bicycle`. I.e. when we combine a more general function with a less general function we get the type of the less general.

Co- *and* contravariant T's in parameters:

```
m<T>(a:T f:(t:T)=>bool): T;  
g: (t:Vehicle)=>bool;  
var r = m(bus, g);
```

Combining what we have seen above, the applicable overloads are those where `T` is a *supertype* of `Vehicle` and a *subtype* of `Vehicle`. Now something interesting happens, though. Whenever an overload is better for the first argument it gets worse for the second. Going by parameter types, *there is no overload!!* Look at the endpoints, `Bus` and `Vehicle` for instance. Clearly `Bus` is better than `Vehicle` for the first parameter. However, almost as clearly `Vehicle=>bool` is better than `Bus=>bool` for the second parameter! So based on parameters, `m<Vehicle>` and `m<Bus>` are equally good, and so are all the overloads in between.

turned
t turned
cable
he most
exist, it is
supertype

r
ve would
ine
be a
ult; it is a

come
ine a

e of Bus
d gets
o best
le for
ne second
che

It is tempting to apply an arbitrary rule here – e.g. “pick the most specific of the T’s”, which in this case would cause us to infer the type of `r` to be `Bus`. That seems right enough in this case. But this method will be wrong. Consider this slight variation:

```
n<T>(a:T f:(t:T=>bool): (t:T=>bool);  
g: (t:Vehicle)=>bool;  
var r = n(bus, g);
```

Everything but the return type (and the name of the method) is the same. But now that arbitrary rule would give us the *least* useful type possible as the return type.

It is time to think deeper about what these overloads really mean. Recall that the overloads represent *one* implementation at runtime. That implementation has *no* access to any type argument we have inferred, and it has *no* access to the compile time types of any of its arguments. It only has access to the arguments *themselves*, including what runtime type information it can glean.

We should therefore feel confident that the return type provided by *any* applicable overload represents a correct (if incomplete) type for what the implementing method would produce on *any* argument satisfying that overload’s static types. Think about that for a second. What that means is that *all applicable overloads are right*. In other words, the actual value returned by the call will have *all the return types that the applicable overloads provide!*

To see this more clearly, take the arguments `bus` and `g` above. Without any unsafe stuff going on, we can do the following:

```
var h: (t:Bus)=>bool = g; // totally safe and legal  
var r2 = m(bus, h);
```

We’ve called the same runtime method with the same objects, but now clearly the result is a `Boolean`. There’s just one applicable overload!

Conversely we can do:

```
var vehicle: Vehicle = bus; // boringly legal  
var r2 = n(vehicle, g);
```

And just as clearly the result is a `(v:Vehicle)=>bool`. Again, there’s just one applicable overload!

In both cases, by *weakening* our static knowledge about an argument, we got the most specific

this case
might as

ry rule

present
happen to
to the

represents
ents
all of the
l the

on I can

us!

ad.

result

in both cases, by weakening our static knowledge about an argument, we get the most specific type possible. That must mean that result type really actually applies! (Unless of course the overloads are not a correct representation of the behavior of the runtime method – but let's deal with that issue another day).

So because all overloads represent the same runtime method, *everything* we know about the runtime value from *every* applicable overload is true. And of course, the combination of all that knowledge is what you guessed it – the intersection of all those return types. The most general common subtype of all the return types.

Co- and contravariant T's in return types:

This is an awesome simple rule to follow and think about: The result type of a method application is the most general common subtype of all the result types of all the applicable overloads.

For generic methods we have to take into account the little snag that there are sometimes infinitely many applicable overloads. So we need some algorithmic handle on this infinity.

For all the results above it so happened that the combined result type was at one of the endpoints of the spectrum - it was the most specific of the return types available. However, there may not be such a specific return type. Let's combine the last two versions of m and n:

```
m<T>(a:T f:(t:T)=>bool): { a:T; f:(t:T)=>bool; }  
g: (t:Vehicle)=>bool;  
var r = m(bus, g);
```

The endpoints of the spectrum are { a:Vehicle; f:(t:Vehicle)=>bool; } and { a:Bus; f:(t:Bus)=>bool; }. However, the combined information tells us more than each of these; namely, that the result is always a { a:Bus; f:(t:Vehicle)=>bool; }. It is in fact amazing that this is the result – it fits so snugly with the arguments we passed in! But we do need a way to find it that terminates within the lifetime of the Universe. I believe there is a straightforward algorithm for this.

The parameters have essentially given rise to an upper and a lower bound for T (Vehicle and Bus in the first case). As we construct the result type we keep track of where T occurs contravariantly and where it occurs covariantly in the declared return type. In the contravariant positions we substitute its upper bound. In the covariant positions we substitute its lower. Done!

Thoughts?

^Mada

... occur
overloads
what kind of

returned
edge is –
of all the

ion is the

nitely

ints of a
ch a most

amely
s the

r this:

us in this
ere it
upper

From: Mads Torgersen
Sent: Tuesday, December 06, 2011 9:47 AM
To: Strada Design Team
Subject: Some thoughts on generic methods

I've been doing some thinking on generic methods, and the below is as much of a summary as for before today's meeting. Looking forward to discuss, and apologies for sending this so close meeting – it had to get written first! ☺

Mads

Generic types seem relatively straightforward. In a structural world such as ours, they are really templating mechanism for types.

The real fun comes in with generic methods. Those cannot in the same way just be thought of as templates – rather they are intricately tied to the “ground level” type system itself. When a type is a generic method, that describes a real-world entity – an actual method present at runtime in an assembly that cannot just be “expanded away.” The underlying runtime reality therefore needs to play a role in how we shape generic methods.

A model of generic methods

How should we think of generic methods? A good place to start is with subtyping. Assume we have the following types:

```
interface A : { f(a: any): any }
interface B : { f(s: string): string }
interface C : { f<T>(t: T): T }
```

What are the subtype relationships between them? There is no subtype relationship between A and B because co- and contravariance of function types work in opposite directions. How about the relationship between C and the other types though? How should we think of the type of C.f and how to compare it to the others?

In C# we don't have such problems. The type system is nominal, and C.f is just a different f than B.f – they have no relationship. When assigned to a delegate a generic method in C# must first

I had time
to the

y just a

as
e has a
object –
role in

have the

A and B

d

n A.f or
be

instantiated with type arguments; it doesn't have a type in and of itself.

The only reasonable model I can think of in Strada is to say that a generic method conceptually bundle of overloads, one for each possible instantiation. The type of $C.f$ is $\{ \langle T \rangle (t: T): T \}$ which is a form of saying $\{ (t: \text{any}): \text{any}; (t: \{\}): \{\}, \dots, (t: \text{number}): \text{number}, (t: \text{string}): \text{string}, \dots \}$; i.e. an infinite number of similarly structured overloads.

This model lets us compare C to the other types. By our usual rules of subtyping, C is a subtype of B and A , because for each overload in each of those, there is a conforming overload in C . Of course, dealing with infinite lists of overloads is interesting territory from an algorithmic perspective, but we'll leave that out: after all these are highly regular sets of overloads induced by a single declaration.

Conceptually I think the model makes a lot of sense. If a set of overloads describes the ways in which a function can be called (and the result types that correspond to each), then this is really an accurate description of what generic methods are for.

This model also means that “inferring the type argument” for a generic method call is an act of overload resolution: which of the infinitely many overloads are *applicable*, and which do we pick? So let's talk some about that.

Overload resolution

Overload resolution is again one of those concepts where we cannot just transfer out experience from C# directly. At runtime there *are* no overloads in Strada – they are purely there to describe the overloads accurately as possible at compile time. So “overload resolution” really has two purposes:

- Deciding whether a given function call is well typed
- Determining the return type of the call expression

Because overload resolution does not have runtime semantic import, there are two kinds of error that really seems unfortunate to have to give as a result:

- Ambiguity: there is no ambiguity at runtime. If two overloads apply equally well, that should be even better than one, and not cause the compiler to give up.
- Getting it wrong: Lack of static information shouldn't be able to lead us down the wrong path by making type assumptions that are patently bad.

Addressing the first one first, what *should* we do if there is more than one applicable overload? It isn't a good way to choose? We could use a *bad* way to choose (i.e. arbitrary), but that would lead to false type assumptions down the road. Or we could somehow “choose all of them” And by that

is a
is a short
nite

of both A
se dealing
figure

which a
rate

overload
s talk

ce from
typing as

rors it

ould be

path,

and there
ead to
t I mean

raise type assumptions down the road. Or we could somehow choose all of them. And by the way, we find a way to “merge” the result types of all the applicable overloads (or at least all the “best” overloads for some meaning of that word) and use that as the type of the call expression.

There are a number of different approaches one could take to this “merging” business:

- Union types – which (confusingly) means types constructed by taking the intersection of sets
- Intersection types – which (just as confusingly) means types constructed by taking the union of member sets
- Something in between – e.g. coming up with a notion of optional members meaning that objects of the type will have those members

Union types are safe and correct, in that they produce a supertype of all the possible return types, therefore adequately describing what they have in common and what can therefore be known about them. However, they are not very useful because that intersection of members is often quite small. The first thing you want to do is assign to the type you *actually* (as a programmer) know it to be.

Intersection types are a lie, but maybe a good lie. They sort of claim that the result has *all* the possible return types at the same time, which is of course preposterous. However, if as a programmer you know what it is *really* supposed to be, it is super useful to be able to directly dot into the right member. For the result to be a subtype of what you know it to be. This might offer the opportunity to get rid of our reverse assignability rule, I may get to that later.

I’ve spent a lot of time in the wilderness of optional members. All I got was a bunch of complex types without much payoff. I think we’d venture there at our own peril, and waste a lot of time getting no value. I think both union type and intersection type approaches are worth investigating.

Applicability

When is a function overload applicable? With our C# hat on we would base that on “implicit convertibility” – i.e. assignability – of arguments to parameter types. There are a number of things to consider though:

- Should we consider expected result types also? After all Java does that, and it could really help to limit the number of applicable candidates nicely. As long as we keep it within the statement, it shouldn’t get *too* complicated.
- What kind of assignability do we consider? Does the argument type have to be a *subtype* of the parameter type, or do we admit overloads that only apply by reverse assignability?

can mean
applicable

member

tion of

some

oes,
about *all*
e trivial.
e.

ossible
ou know
ers, and
t rid of

ity
ng little

ngs to

y help
ent, it

of the

I won't get into too many pro's and cons now. One thing I will dive into: with the right combination of settings we might get rid of reverse assignability from the language! How so? The motivating scenario for reverse assignability is for factories that produce a number of different kinds of things to be able to have their return value (typed by a common supertype) directly assigned to one of those products. However, if we describe such a factory method as a generic method:

```
make<T>(recipe: string): T
```

And we allow overload resolution (and hence type inference) to take expected type into account, then we can call this as:

```
Result: MyKindOfProduct = make("MyKindOfProduct");
```

`MyKindOfProduct` would become an upper bound on `T` in the call, and hence only the overload where `T` is a subtype of `MyKindOfProduct` would apply. Similarly you could consider

```
make("MyKindOfProduct").Foo(7);
```

to be legal, imposing `{ Foo(a: int): any }` as an upper bound on `T`. In general, the expectations on the return type of "make" can probably be summarized as a structural type.

There are guaranteed to be subtleties here, but it is an interesting idea to pursue.

Constraints

One simplification we've considered is to not have constraints on type parameters. If we go with the "merge" model for the return types of applicable overloads, this might not be the best option. When you call a generic method, the argument types (and also perhaps expected result types) will introduce upper and lower bounds on type parameters. To the extent that the result type is itself generic, it might be better (at least more precise) to capture those upper and lower bounds on the type parameters and use them to constrain that result type.

Lambda arguments

In C# we infer through lambda arguments by "pushing in" known parameter types and seeing what comes out the other end.

There is an alternative possible approach in Strada: We could find the type of the function expression compositionally, without considering its context. We could do that by letting each unspecified parameter type of the lambda give rise to a fresh type parameter on the (generic) function type.

tion of
scenario
e able to
uct types.

nt, then

s where T

n the

th a
When we
uce both
, it would
ters of

what we

ession

e that we

are inferring, and then collecting constraints on that type parameter based on how the corresponding parameter is used in the lambda.

Mads

onding